

CPSC 313: Computer Hardware and Operating Systems  
Assignment #4, due Thursday, April 6th, 2017 at 23:59

See the edX assignment page for additional administrative details.

## Introduction and Objectives

This assignment is all about file systems. While completing this assignment:

1. You will gain additional practical experience writing code in the C programming language.
2. You will gain a better understanding of the implementation of at least one file system (specifically, the MS-DOS file system that is used by most digital cameras and MP3 players).

The assignment is divided into three parts. The division is meant to provide guidelines to help you allocate time when you are working on the assignment, and to help you make steady progress on it instead of leaving it to the last minute. Although the assignment is broken into three parts you are not required to handin the intermediate parts and will handin only the completed assignment. However, you are strongly encouraged to push things to Stash/Bitbucket on a regular basis. Parts of the assignment can be challenging so you should get started early. If you want you can work with one other person. Groups of three or more people are not permitted.

The program you are to write, named **fatinfo** is to be written from scratch using C except for some code provided to simulate a disk drive and to get you started. The program must compile and run on the undergraduate Linux machines provided by the department. **Programs that do not compile will be awarded 0.** If you are using your own machine, make sure to test that it compiles and runs as expected before handing it in. Keep in mind that the department machines are all running a 64 bit version of Linux as a result there are lots of things that could differ between your machine and the department environment. You are not allowed to use C++ to complete this assignment. You are only allowed to use functions from the standard C library. In particular, this means that you are *not* allowed to use a library of functions that deal with FAT file systems in any way, regardless of its source.

When writing this program you are to use good software design practices for the naming of functions, variables, and the organization of functions into separate files. In particular this means that you are not to put all of your functions into a single file. Instead, put related functions into the same file. For example, putting functions dealing with block I/O together, and then putting directory functions together would make a lot of sense.

Unlike the previous assignments you are being provided with very little initial code. The provided repository contains a suggested collection of files to place your code, function prototypes, structure definitions etc. into, along with a Makefile. These files are just suggestions and you are free to delete, rename, or add to the collection of files that comprise

this assignment, however you must keep your git repository upto date and it must always be the case that the following make commands work:

- *make clean* - this will always clean out the .o files and the executables
- *make depend* - this will rewrite the makefile to build a proper list of dependencies. As long as you keep the SRC line in the Makefile upto date this should work. If you add or delete any file from the src line be sure to run *make depend* and then commit the updated Makefile.
- *make* or *make fatinfo* - this will cause the fatinfo program to be built. NOTE your program must build and compile without producing any warnings. You are not permitted to add compiler directives or directives into the C code to suppress warnings. However, you may get warnings when you do a *make depend* and that is acceptable.
- The file disk.c contains a collection of routines to simulate accessing a disk. The only way you are allowed to access an image of a disk file is through these routines. You are not allowed to modify these functions without the permission of course staff.
- The file *fatinfo.c* has some starter code in it. You are free to use that code or write your own.

## Part 1

You should aim at completing this part of the assignment by March 25 to ensure you have enough time to finish the assignment or get help if needed. The goal of this part is to help you become comfortable using random access techniques to access the data in a file and then interpreting some of the files contents. On edX, there are several files that contain images of various FAT file systems. The FAT file system comes in various formats and your code is required to work with both FAT12 and FAT32 versions. The primary difference is in the way information is represented in the FAT table and the fact that FAT32 uses 32 bit addresses while FAT12 uses 12 bit addresses.

The task for this first part is to write code to output the following information about this file system:

- its sector size.
- its cluster size in sectors.
- the number of entries in its root directory. - This is only meaningful for FAT12. For FAT32 the root directory is like all other directories and does not have a fixed size whereas the FAT12 root directories have their size fixed when the file system is initialized. Consequently for FAT32 you will print 0.
- the number of sectors per file allocation table.

- the number of reserved sectors on the disk.
- the number of hidden sectors on the disk.
- the sector number of the first copy of the file allocation table.
- the sector number of the first sector of the root directory. - The way to determine this varies between FAT12 and FAT32.
- the sector number of the first sector of the first usable data cluster. Keep in mind that in both systems the first 2 clusters (0 and 1) are "reserved" and are not used. (The FAT table will have entries for these clusters but the values in the FAT table have a special meaning.) Also in FAT12 cluster 2 starts right after the root directory whereas in FAT32 it starts right after the FATs.
- a function in the fatinfo.c file takes a structure with this information in it and prints the information in the required order and format. **DO NOT DEVIATE FROM THIS FORMAT AND ORDER**

All of this information can be computed from the information stored in the first (BOOT) sector of the file system. You may assume that if bytes 17-18 are 0 then you are dealing with a FAT32 system otherwise you are dealing with a FAT12. Other than figuring out some of the basic information about the file system the other major area these two systems differ in is how the FAT is used. Many of the details about the MS-DOS file systems can be found here:

<http://www.win.tue.nl/~aeb/linux/fs/fat/fat-1.html>

(Additional information is available in the section of edX devoted to this assignment.) One of the sentences on this page is somewhat confusing. Under the heading "FAT12", you will find the statement:

Since 12 bits is not an integral number of bytes, we have to specify how these are arranged. Two FAT12 entries are stored into three bytes; if these bytes are  $uv$ ,  $wx$ ,  $yz$  then the entries are  $xuv$  and  $yzw$ .

Here is a simpler interpretation of the second sentence: Suppose you have the bytes  $b_1$ ,  $b_2$  and  $b_3$  in that order. You can extract the two FAT12 entries by treating these three bytes as a single little-endian 24 bit unsigned integer  $x$ . To get the FAT12 entry that occupies the low-order 12 bits *and*  $x$  with 0x0FFF. To get the FAT12 entry that occupies bits 12 to 23 bits shift the  $x$  value to the right by 12 so that bits 12 to 23 are now bits 0 to 11. Unless you are 100% sure that the top order 8 bits were 0, *and* this value with 0x0FFF to ensure that only the 12 bits of interest are captured.

Once you have extracted the information from the first sector you are to print the information out in the exact format shown in the sample output.

When writing your code for part 1, be sure to make it general enough to use again for parts 2 and 3. Note that you are not allowed to use any I/O functions to read the contents of the disk images other than the ones provide.

To modularize your code you might find it useful to define a `struct` like `filesystem_info` that contains most of the information about the file system that part 1 asks you to determine and to encourage you to do this such a structure has been provided. It is strongly recommend that you write functions to:

- Read a single cluster (In the DOS world the term cluster corresponds to the term block that we have used in class) into a buffer in memory.
- Extract an unsigned integer value from two (or three, or four) specific bytes treated as a little-endian value.

Your test program must take a single argument, the name of the file containing the disk image of the file system it is supposed to interpret.

## Part 2

You should aim at completing this part of the assignment by Thursday March 30. The goal of this part is to be able to deal with directory entries, and extract information about the files a directory contains. More specifically, you should add to your code from part 1 so it can output the following information about each file in the root directory:

- Its name, including the extension if one is present.
- Whether or not this file is a directory.
- The number of the first cluster containing its data.
- Its size (note that directories have a size of 0).

This part of the assignment should be relatively straight forward once you have completed part 1. Note that the format for the directory entries is exactly the same for FAT12 and FAT32. One thing to be aware of is that if the parent of a subdirectory is the root directory then the cluster number for that directory is sometimes given as 0, so you will have to be careful about that in the next section. (Recall that the cluster 0 and 1 are reserved so using it is a special value is OK.)

## Part 3

To complete this part, the program is to print out the information from part 2 about every file in the file system, **including files in subdirectories**. The name printed for each file is to include its full path name (for instance, `ASS2/MCHEME.TXT`). The program is also to print

out the list of clusters used for each file using the same format as provided in the sample output that can be found in this assignment section on edX.

This part of the assignment will likely be the most challenging because it deals directly with the File Allocation Table. To work with the FAT it is recommended that you write functions that:

- determine if a given cluster number is the last cluster in its file.
- retrieve the next cluster number in a file, given the current cluster number.

When printing the information observe the following requirements:

1. Use the exact formatting shown in the samples paying particular attention to column order and and spacing.
2. Within a directory, print the entries in the order that they appear.
3. If an entry is a directory, print the information about it and then enter that directory and recursively print/traverse that directory and any of its sub-directories. (Basically you are doing a depth-first search of the filesystem.)
4. The clusters associated with a file or directory are to be printed in the order that those clusters would have to be accessed if reading the file, in order, from the first byte to the last byte.
5. A collection of contiguous clusters are to be printed as a range. For example 1, 2, 3 , 4 is to be printed as 1-4 (note - no spaces). Something like 1, 2, 3, 4, 8, 9, 7, 20, 21, 23 is to be printed as: 1-4,8-9,7,20-23.
6. Print all the information for a file/directory on one line, regardless of how long the line is.
7. Make sure that your program signals that no more clusters are associated with a file by printing [END].

## Deliverables

You are to submit the assignment using git, just like you have for all previous assignments in this course. The files to submit for this part are:

1. All of your C source and include files.
2. A makefile such that when make is typed in the directory containing your solution the program **fatinfo** is built. This program takes a single argument, the name of the file containing the sample file system the program is to decode. Your code must compile by typing make and there must be no errors or warnings. You are not allowed to change the gcc options in the makefile to alter how and when any warning messages are printed.

3. The filled in version of `coverpage.txt`. Follow the directions in the `coverpage.txt` file on how to fill-in and submit the file.